

Highest-performance Stream Processing

Oleg Kiselyov

Tohoku University
Japan
oleg@okmij.org

Aggelos Biboudis

Switzerland
biboudis@gmail.com

Tomoaki Kobayashi

Tohoku University
Japan
tomoaki.kobayashi.t3@dc.tohoku.ac.jp

Nick Palladinis

Nessos IT
Greece
npal@nessos.gr

Abstract

We present the stream processing library that achieves the highest performance of existing OCaml streaming libraries, attaining the speed and memory efficiency of hand-written state machines. It supports finite and infinite streams with the familiar declarative interface, of *any* combination of map, filter, take(while), drop(while), zip, flatmap combinators and tupling. Experienced users may use the lower-level interface of stateful streams and implement accumulating maps, compression and windowing. The library is based on assured code generation (at present, of OCaml and C) and guarantees in all cases *complete fusion*.

1 Summary

Strymonas is a DSL that generates high-performance single-core stream processing code from declarative descriptions of stream pipelines and user actions – something like Yacc. Unlike (ocaml)yacc, strymonas is an embedded DSL. Therefore, it integrates as is with the existing OCaml code and tools. Any typing or pipeline mis-assembling errors are reported immediately (even during editing).

Strymonas statically guarantees *complete fusion*: if each operation in a pipeline individually runs without any function calls and memory allocations, the entire streaming pipeline runs without calls and allocations. Thus strymonas per se introduces not even constant-size intermediary data structures. Complete fusion is mainly the space guarantee: the ability to run the processing loop without any GC or even stack allocations. Still, avoiding closures and the repeated construction/disposal of tuples, option values, etc. notably improves performance, in our experience.

The present strymonas is the completely re-written and the much extended and improved version of the library described [Kiselyov et al. 2017]. The main differences are:

- Generated code is not only OCaml but also C. The latter needs no OCaml runtime and can be automatically vectorized.
- The core of strymonas – stream representation and fusion – is now pure OCaml, with code generation relegated to a backend. MetaOCaml is needed only for the OCaml-code-generation backend – but not for the C backend. Users may develop their own backends by implementing a specific signature.
- Stream fusion is now achieved in all cases. In [Kiselyov et al. 2017], complicated zipping pipelines did not fuse completely.
- To guarantee the complete fusion in all cases, including complicated zip pipelines, the library was completely re-written, using a different fusion approach based on normalization-by-evaluation.

- The library is structured as the low-level core interface of stateful streams upon which the familiar declarative combinators are implemented. The core interface can be used directly, to write stream processing that depends not only on the current but past stream elements – such as signal-processing filtering, windowing, compression. One may freely mix the two interfaces in the same pipeline.
- The low-level interface supports streams not only over base types (integers, strings, floats) but also over tuples, records and even abstract data types. The example of the latter is windowing.

1.1 Backends

Semantic actions – stream mapping, filtering, accumulating, etc. functions – are expressed in another embedded DSL, called a backend. Strymonas currently provides two implementations of the DSL (with the identical interface): OCaml and C.¹ We are considering WASM and LLVM IR backends.

The OCaml backend is based on MetaOCaml and extends the common backend interface to permit arbitrary OCaml code (enclosed in MetaOCaml brackets) as semantic actions.² On the other hand, the C backend is pure OCaml: MetaOCaml installation is *not* needed. No other dependencies are needed either. The C backend uses the tagless-final-based approach described in [Kiselyov 2022]. Thanks to tagless-final, the backends are extensible. Either backend generates code that is statically assured to compile without errors or warnings.

2 A Taste of Strymonas

Strymonas may be thought of as Yacc for stream processing – but embedded rather than standalone. Here is the simplest example (for clarity, we explicitly write type annotations although none are needed):

```
let ex1 : int cstream = iota C.(int 1) > map C.(fun e → e * e)
```

Like Yacc, strymonas uses two languages: one to describe the structure of the stream pipeline, and the other to specify the semantic actions such as mapping transformations, etc. Since strymonas is an embedded DSL, both languages are represented by OCaml functions (combinators), but from two different namespaces (signatures). Stream structure combinators such as *iota* and *map* produce, consume, or transform values of the type α cstream, where α is a base

¹Actually, there are four backends: basic OCaml, basic C, and the backends obtained by applying a partial-evaluation functor to any backend.

²If such an extensibility is not needed, the OCaml backend can be implemented in pure OCaml: that is, without MetaOCaml.

type.³ The `iota` combinator produces an infinite stream of numbers starting with the given one; `map` should be self-explanatory. The combinators are typically composed via `▷`: the right-to-left application operator. Semantic actions (i.e., the arguments of stream combinators) are described via backend combinators, which build values of the abstract type α `cde` representing the target code: C, OCaml, etc.⁴ We shall assume one such backend in scope, as the module named `C`.

The `ex1` pipeline is the infinite stream of natural numbers transformed by squaring each item. The pipeline is completed by terminating it by an int cstream consumer, such as `fold`, or its instance `sum` below (strymonas already provides this combinator, under the name `sum_int`):

```
let sum : int cstream → int cde = fold C.(+) C.(int 0)
let ex2 : int cde =
  ex1 ▷ filter C.(fun e → e mod (int 17) > int 7)
  ▷ take C.(int 10) ▷ sum
```

We added to `ex1` two more transformations, to retain only those items whose remainder mod 17 exceeds 7, and take first 10 such items. The resulting stream becomes finite and can be meaningfully summed up. Here `ex2` is the int `cde` value representing the integer-producing target code for the complete pipeline, including stream generation, folding, and the user actions of squaring, etc. The backend that realizes α `cde` as OCaml code lets us pretty-print this code:

```
let v_1 = ref 0 in let v_2 = ref 10 in let v_3 = ref 1 in
while (! v_2) > 0 do
  let t_4 = ! v_3 in incr v_3;
  let t_5 = t_4 * t_4 in
  if (t_5 mod 17) > 7 then (decr v_2; v_1 := ! v_1 + t_5)
done;
! v_1
```

The code can also be saved into a file, compiled, put into a library – or it can be run right away: dynamically linked into the program that generated it and invoked. With the C back-end, the resulting code is (some newlines are removed for compactness):

```
int fn()
{ int v_1 = 0; int v_2 = 10; int v_3 = 1;
  while (v_2 > 0)
  { int t_4; int t_5;
    t_4 = v_3;
    v_3++; t_5 = t_4 * t_4;
    if ((t_5 % 17) > 7) { v_2--; v_1 = v_1 + t_5; }
  }
  return v_1;}
```

This is what a competent programmer would have written by hand. Although the pipeline is purely declarative, with first-class (the argument of `filter`) and higher-order functions, the generated code is imperative and has no function calls. The main loops runs with no GC, even in OCaml.

³In a lower-level strymonas interface, stream elements are not restricted to base types.

⁴Thus α `cde` is an abstract, backend-independent representation of the generated code, which may not even be an OCaml code – and hence different from MetaOCaml’s α code. MetaOCaml may not be needed, depending on the chosen backend.

The second example is the pipeline to compute the dot-product of two arrays:

```
let ex_dot (arr1:int array cde,arr2:int array cde) : int cde =
  zip_with C.( * ) (of_arr arr1) (of_arr arr2) ▷ sum
```

The combinator `of_arr` creates a finite stream whose contents is the given target language array. With the OCaml backend we generate:

```
fun (arg1_24, arg2_25) →
  let t_26 = (Array.length arg2_25) - 1 in
  let t_27 = (Array.length arg1_24) - 1 in
  let v_28 = ref 0 in
  for i_29 = 0 to if t_27 < t_26 then t_27 else t_26 do
    let el_30 = Array.get arg1_24 i_29 in
    let el_31 = Array.get arg2_25 i_29 in
    v_28 := (! v_28) + (el_30 * el_31)
  done;
  ! v_28
```

Incidentally, `zip_with` `f` appearing in `ex_dot` is defined in `strymonas` as `zip ▷ map (fun (x,y) → f x y)` where `▷` is left-to-right function composition. A naive implementation would construct a tuple in `zip`, to be deconstructed in the subsequent mapping. The `strymonas`-generated code however clearly has no tuples.

3 Evaluation

We evaluated `strymonas` on the set of micro-benchmarks borrowed from [Kiselyov et al. 2017], to which we added `ZipFilterFilter`, `ZipFlatMapFlatMap` (that is, zipping of two streams each containing a filter, resp., flatmap operation) and `runLengthDecoding` benchmarks. Fig. 1 presents the results. Baseline is the hand-written, hand-fused imperative (state-machine) code for the entire pipeline, with no closures or thunks: what a competent programmer would write to achieve best performance. Streaming⁵ is one of the fastest streaming libraries in OCaml; it does not support all the benchmark cases however.

We have also implemented larger applications/benchmarks, such as FM Radio reception. The generated C code, which compiles to fully vectorized machine code, has the performance to sustain real-time reception.

References

- Oleg Kiselyov. Generating C. In Michael Hanus and Atsushi Igarashi, editors, *Functional and Logic Programming*, volume 13215 of *Lecture Notes in Computer Science*, pages 75–93. Springer International Publishing, 2022. doi: 10.1007/978-3-030-99461-7_5.
- Oleg Kiselyov, Aggelos Biboudis, Nick Palladinis, and Yannis Smaragdakis. Stream fusion, to completeness. In *POPL ’17: Conference Record of the Annual ACM Symposium on Principles of Programming Languages*, pages 285–299, New York, January 2017. ACM Press. ISBN 978-1-4503-4660-3. doi: 10.1145/3009837.

⁵<https://github.com/odis-labs/streaming>




Figure 1. Benchmarking against the baseline and streaming: the running time in milliseconds per iteration (avg. of 20, with mean-error bars shown). Shorter bars are better. To better show the details, the figure is truncated: the per-iteration running time of streaming on mapsMegamorphic is 3.7 sec, on filtersMegamorphic is 2.4 sec. The evaluation platform is 1.8 GHz dualcore Intel Core i5, 8 GB DDR3 main memory, macOS Big Sur 11.6.